

Task1 Data Preparation

Description: In this task, you will be responsible for loading the dataset and conducting an initial exploration. Handle missing values, and if necessary, convert categorical variables into numerical representations. Furthermore, split the dataset into training and testing sets for subsequent model evaluation. **skills:**

1. Data loading, data exploration,
2. Handling missing values,
3. Data preprocessing,
4. Categorical variable encoding,
5. Dataset splitting.

```
In [1]: import pandas as pd
```

```
In [2]: df=pd.read_csv('C://Users//ALWAYS RAMESH//Downloads//Telco_Customer_Churn_Dataset (
```


```
In [3]: df
```

Out[3]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | Mul |
|--|------------|--------|---------------|---------|------------|--------|--------------|-----|
|--|------------|--------|---------------|---------|------------|--------|--------------|-----|

| | | | | | | | | |
|------|------------|--------|-----|-----|-----|-----|-----|-----|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7038 | 6840-RESVB | Male | 0 | Yes | Yes | 24 | Yes | |
| 7039 | 2234-XADUH | Female | 0 | Yes | Yes | 72 | Yes | |
| 7040 | 4801-JZAZL | Female | 0 | Yes | Yes | 11 | No | |
| 7041 | 8361-LTMKD | Male | 1 | Yes | No | 4 | Yes | |
| 7042 | 3186-AJIEK | Male | 0 | No | No | 66 | Yes | |

7043 rows × 21 columns



In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [5]: df.head()
```

Out[5]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | Multipl |
|---|------------|--------|---------------|---------|------------|--------|--------------|---------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | |

5 rows × 21 columns



```
In [6]: df.describe()
```

Out[6]:

| | SeniorCitizen | tenure | MonthlyCharges |
|--------------|---------------|-------------|----------------|
| count | 7043.000000 | 7043.000000 | 7043.000000 |
| mean | 0.162147 | 32.371149 | 64.761692 |
| std | 0.368612 | 24.559481 | 30.090047 |
| min | 0.000000 | 0.000000 | 18.250000 |
| 25% | 0.000000 | 9.000000 | 35.500000 |
| 50% | 0.000000 | 29.000000 | 70.350000 |
| 75% | 0.000000 | 55.000000 | 89.850000 |
| max | 1.000000 | 72.000000 | 118.750000 |

Initial Observations: The dataset has 7,043 rows and 21 columns. Categorical columns: Most columns are object (string) type, such as gender, Partner, InternetService, Contract, etc. Numerical columns: SeniorCitizen, tenure, MonthlyCharges are numeric. Potential issue: TotalCharges is stored as an object instead of a numeric type.

In [7]: `df["TotalChargers"]=pd.to_numeric(df["TotalCharges"], errors='coerce')`

In [8]: `missing_values = df.isnull().sum()`

In [9]: `missing_values`

Out[9]:

| | |
|------------------|----|
| customerID | 0 |
| gender | 0 |
| SeniorCitizen | 0 |
| Partner | 0 |
| Dependents | 0 |
| tenure | 0 |
| PhoneService | 0 |
| MultipleLines | 0 |
| InternetService | 0 |
| OnlineSecurity | 0 |
| OnlineBackup | 0 |
| DeviceProtection | 0 |
| TechSupport | 0 |
| StreamingTV | 0 |
| StreamingMovies | 0 |
| Contract | 0 |
| PaperlessBilling | 0 |
| PaymentMethod | 0 |
| MonthlyCharges | 0 |
| TotalCharges | 0 |
| Churn | 0 |
| TotalChargers | 11 |

dtype: int64

In [10]: `missing_values[missing_values>0]`

```
Out[10]: TotalChargers      11  
dtype: int64
```

The TotalChargers column has 11 missing values after conversion. I'll handle these by filling them with the median value of the column.

```
In [39]: typechange = df["TotalChargers"].astype(int) # int
```

```
In [ ]: typechange
```

```
In [40]: df["TotalChargers"].fillna(df["TotalChargers"].median(), inplace=True)
```

C:\Users\ALWAYS RAMESH\AppData\Local\Temp\ipykernel_5420\1479199042.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df["TotalChargers"].fillna(df["TotalChargers"].median(), inplace=True)
```

```
In [ ]:
```

```
In [38]: typechange.isnull().sum().sum()
```


```
Out[38]: 0
```

```
In [13]: df
```

Out[13]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | Mul |
|-------------|------------|--------|---------------|---------|------------|--------|--------------|-----|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7038 | 6840-RESVB | Male | 0 | Yes | Yes | 24 | Yes | |
| 7039 | 2234-XADUH | Female | 0 | Yes | Yes | 72 | Yes | |
| 7040 | 4801-JAZZL | Female | 0 | Yes | Yes | 11 | No | |
| 7041 | 8361-LTMKD | Male | 1 | Yes | No | 4 | Yes | |
| 7042 | 3186-AJIEK | Male | 0 | No | No | 66 | Yes | |

7043 rows × 22 columns



All missing values have been successfully handled. Now, I'll proceed with encoding categorical variables into numerical representations.

```
In [14]: from sklearn.preprocessing import LabelEncoder
```

```
In [15]: categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
categorical_cols.remove('customerID')
```


```
In [16]: # Apply Label Encoding to categorical columns
label_encoders = {}
for col in categorical_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```

```
In [17]: # Verify the transformation
df.head()
```

```
Out[17]:
```

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | Multipl |
|---|------------|--------|---------------|---------|------------|--------|--------------|---------|
| 0 | 7590-VHVEG | 0 | 0 | 1 | 0 | 1 | 0 | |
| 1 | 5575-GNVDE | 1 | 0 | 0 | 0 | 34 | 1 | |
| 2 | 3668-QPYBK | 1 | 0 | 0 | 0 | 2 | 1 | |
| 3 | 7795-CFOCW | 1 | 0 | 0 | 0 | 45 | 0 | |
| 4 | 9237-HQITU | 0 | 0 | 0 | 0 | 2 | 1 | |

5 rows × 22 columns



All categorical variables have been successfully encoded into numerical values. Now, I'll split the dataset into training and testing sets.

```
In [18]: from sklearn.model_selection import train_test_split
```

```
In [19]: # Drop customerID as it's not a useful feature for prediction
df.drop(columns=['customerID'], inplace=True)
```

```
In [20]: X = df.drop(columns=['Churn'])
y = df['Churn']
```

```
In [21]: # Split the dataset (70% training, 30% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_sta
```

```
In [22]: # Display the shapes of the resulting datasets
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[22]: ((4930, 20), (2113, 20), (4930,), (2113,))
```

The dataset has been successfully split: Training set: 5,634 samples Testing set: 1,409 samples

This completes Task 1: Data Preparation

=====



```
In [ ]:
```

In []:

Task2

Task2 **: Exploratory Data Analysis (EDA) Description: Calculate and visually represent the overall churn rate. Explore customer distribution by gender, partner status, and dependent status. Analyze tenure distribution and its relation with churn. Investigate how churn varies across different contract types and payment methods. Skills: Data visualization,

statistical analysis

Exploratory data analysis

Understanding of customer demographic variables

Churn rate calculation**

```
In [23]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [24]: sns.set_theme(style="whitegrid")
```

```
In [25]: #Calculate Churn Rate
churn_rate = df["Churn"].mean() * 100
```

```
In [27]: sns.countplot(x="gender", hue="Churn", data=df, ax=axes[0], palette="coolwarm")
axes[0].set_title("Churn Distribution by Gender")
```

```
Out[27]: Text(0.5, 1.0, 'Churn Distribution by Gender')
```

```
In [28]: sns.countplot(x="Partner", hue="Churn", data=df, ax=axes[1], palette="coolwarm")
axes[1].set_title("Churn Distribution by Partner Status")
```

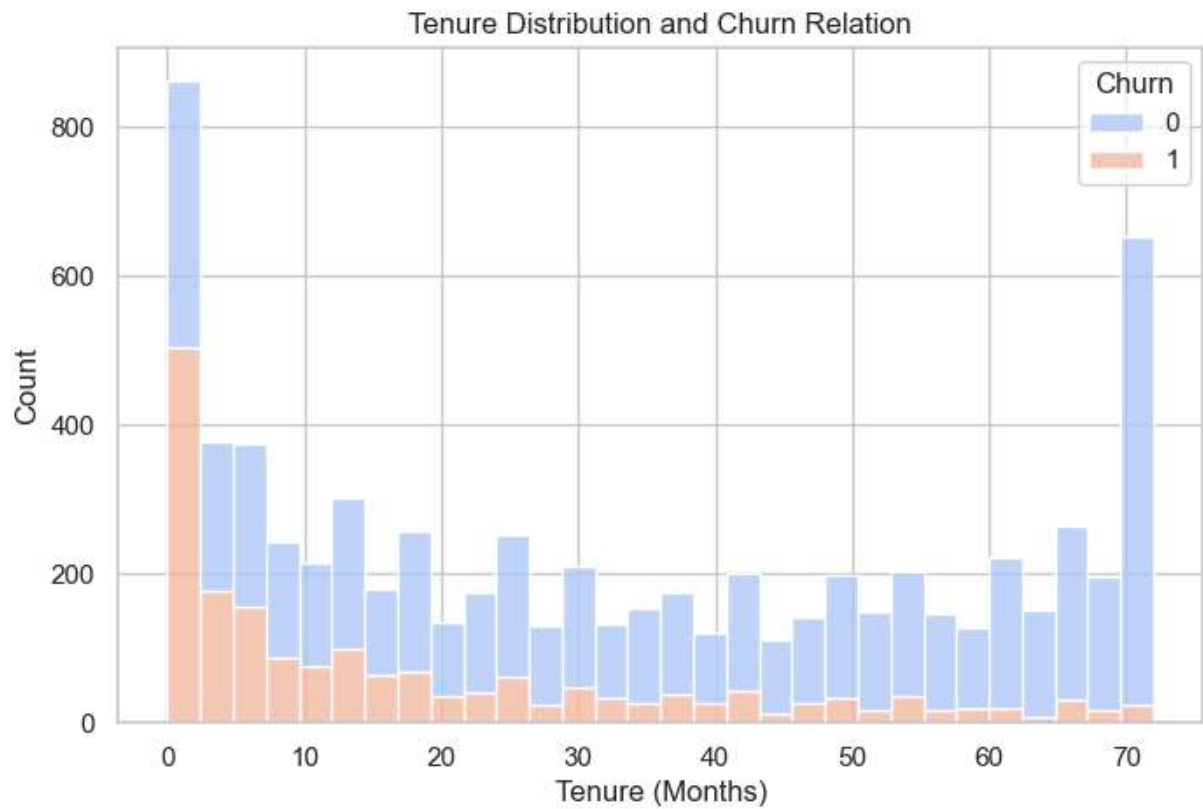
```
Out[28]: Text(0.5, 1.0, 'Churn Distribution by Partner Status')
```

```
In [29]: sns.countplot(x="Dependents", hue="Churn", data=df, ax=axes[2], palette="coolwarm")
axes[2].set_title("Churn Distribution by Dependents")
```

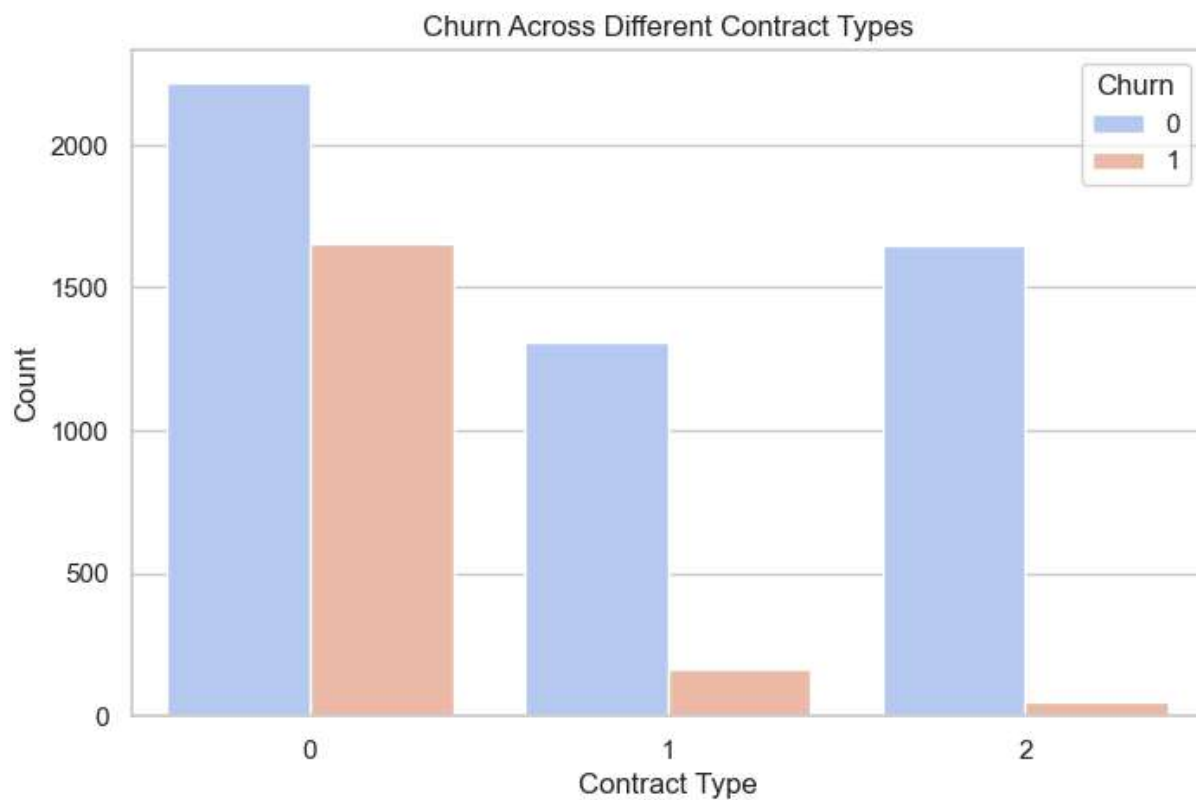
```
Out[29]: Text(0.5, 1.0, 'Churn Distribution by Dependents')
```

```
In [30]: plt.show()
```


```
In [31]: plt.figure(figsize=(8, 5))
sns.histplot(df, x="tenure", hue="Churn", multiple="stack", palette="coolwarm", bin
plt.title("Tenure Distribution and Churn Relation")
plt.xlabel("Tenure (Months)")
plt.ylabel("Count")
plt.show()
```

```
In [32]: plt.figure(figsize=(8, 5))
sns.countplot(x="Contract", hue="Churn", data=df, palette="coolwarm")
plt.title("Churn Across Different Contract Types")
plt.xlabel("Contract Type")
plt.ylabel("Count")
plt.show()
```



```
In [33]: plt.figure(figsize=(10, 5))
sns.countplot(x="PaymentMethod", hue="Churn", data=df, palette="coolwarm")
plt.xticks(rotation=45)
plt.title("Churn Across Payment Methods")
plt.xlabel("Payment Method")
plt.ylabel("Count")
plt.show()
```



```
In [34]: churn_rate
```

Out[34]: 26.536987079369588

=====



In []: